
TS2VecAR - Adding autoregressive temporal contrasting to an universal representation of a time series

Constantin von Crailsheim

Department of Statistics

Ludwig-Maximilians-Universität München c.crailsheim@campus.lmu.de

Abstract

While classification of time series data is crucial in various domains, there are only few labelled datasets available. Thus, self-supervised learning provides significant added value to learn the inherent structure of a time series before fine-tuning a classifier. This investigation builds on TS2Vec (Yue et al., 2021), which learns a robust contextual representations of a time series that can be used for several downstream tasks. Inspired by Eldele et al. (2022), we add a cross-prediction task of future time stamps using information about previous time stamps as summarised by an autoregressive model. This allows the extended TS2VecAR model to learn more about the structure of the time series. The model outperforms the model on several benchmark datasets, with better average performance on Human Activity Recognition (HAR) datasets. The source code is available at <https://github.com/constantin-crailsheim/TS2VecAR>.

1 Introduction and related work

Time series data play an integral role in science and industry in various fields such as medicine, finance and manufacturing. However, Eldele et al. (2022) pointed out that human annotation is very challenging since time series patterns are not easily recognisable by humans, hence only few time series data have been labelled (Ching et al., 2018). This demonstrates the need for self-supervised learning methods which can learn the structure of an unlabelled time series in a pretext task and thus require only few labelled instances to fine-tune a classifier.

Self-supervised representation learning has been particularly popular in computer vision using a contrastive loss function. In contrastive methods, invariant representations of the initial data are learned by inducing similar latent representations for the same instances in a different augmented context and dissimilar latent representations for different instances. Bachman et al. (2019) maximize the mutual information between features in multiple views to induce the algorithm to learn higher-level properties of the data. Chen et al. (2020) propose a simplified contrastive learning framework and highlight the role of data augmentation configurations.

Representation learning for time series has recently gained momentum. Franceschi et al. (2019) obtained a generic representation by sampling positives as random subseries and feeding these through a dilated convolutional encoder, whereby the latent representations are evaluated with triplet loss. Mohsenvand et al. (2020) extended the SimCLR framework (Chen et al., 2020) to perform a classification task on EEG time series data. Tonekaboni et al. (2021) propose a contrastive learning framework for non-stationary time series, where the distribution of local signals should be distinguishable. Oord et al. (2018) use an autoregressive model to predict future instances in

the latent space to induce representation that captures relevant information for predicting future instances.

While previous work has yielded latent representations that satisfy subseries consistency (Franceschi et al., 2019) and temporal consistency (Tonekaboni et al., 2021), Yue et al. (2021) argue that these strong assumptions may be violated in the presence of level shifts and anomalies. Thus, they propose contextual consistency, which simply "treats the representations at the same timestamp in two augmented contexts as a positive pair" (Yue et al., 2021, p. 8982). Since they showed very strong results in several tasks, we used TS2Vec as the underlying framework for my model implementation. However, TS2Vec only evaluates the quality of the representation of each timestamp in a rather isolated way and does not learn the structure of the time series in an autoregressive sense. Therefore, we integrated the cross-prediction task proposed by Eldele et al. (2022) in their TS-TCC model into the TS2Vec framework, which uses a context vector summarising a sample of consecutive latent representations with a Transformer (Vaswani et al., 2017) as autoregressive model to predict future latent representations. With this extension, which we will refer to as TS2VecAR, we aim to derive robust contextual representations that also capture relevant information about future timestamps.

2 Method

As shown below, TS2VecAR embeds the temporal contrasting module of TS-TCC into the initial implementation of TS2Vec.

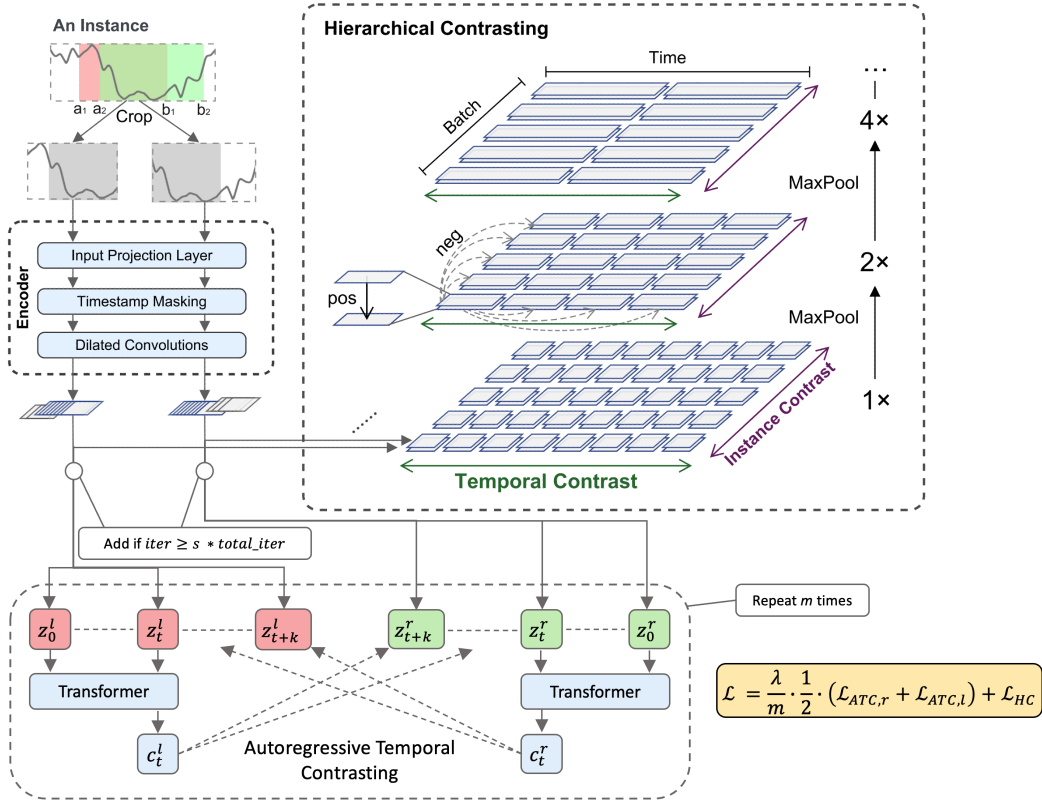


Figure 1: Structure of model (TS2Vec part copied from Yue et al. (2021) and ATC part is own illustration based on Eldele et al. (2022))

In TS2Vec, two overlapping windows of a (multivariate) time series are sampled, where the left window is defined by $[a_1, b_1]$ and the right window by $[a_2, b_2]$. The following relationship has to hold for the cut-off points: $0 < a_1 \leq a_2 \leq b_1 \leq b_2 \leq T$. Each window is encoded separately by feeding

it through an input projection layer, timestamp masking and dilated convolutions. The similarity of the overlapping part (i.e. $[a_2, b_1]$) of the encoded sequences in two different contexts is evaluated using hierarchical contrasting, applying temporal and instance contrasting iteratively. Temporal contrasting applies a contrastive loss over all time stamps of the same time series, where the same time stamps encoded in two representations are treated as positives and different time stamps as negatives. Instance contrasting, meanwhile, treats the two representations from different contexts of the same time series as positives and representations of different time series in the batch as negatives. By employing this two-fold loss function, the authors aim to achieve contextual consistency, which induces more robust learned representations. See Yue et al. (2021) for more details (necessary?).

As outlined above, temporal contrasting considers time stamps rather isolated. Thus, we integrated the idea by Eldele et al. (2022) to use an autoregressive model, which summarizes the latent representation into a context vector, into the TS2Vec framework. To be precise, all latent representations up to a randomly sampled timestamp t (i.e. $z_{<t}$) are passed to a Transformer, which summarizes them into a context vector c_t . This context vector is used to predict the next k latent representations that were encoded from the other window (cross-prediction task). The prediction is evaluated using a contrastive loss over all time stamps in the prediction window. See Eldele et al. (2022) for more technical details.

The main adaptations are the following:

- While Eldele et al. (2022) use so-called strong and weak augmentations to induce latent representations of the time series in two different contexts, here two sampled and overlapping windows are used according to the TS2Vec model structure.
- The autoregressive model may not be included until later iterations of the optimization, allowing the original TS2Vec component to induce reasonable latent representations that are refined by an autoregressive component at a later stage.
- Since the context vector is only used to predict a limited number of k consecutive latent representations at each iteration, the procedure of sampling t and cross-predicting using a context vector can be repeated m times to exploit a larger fraction of the time series.
- The final loss function is the sum of the hierarchical contrastive loss and the mean of the two autoregressive temporal contrastive losses, which were derived by cross predicting the latent representations encoded from the left and right windows.
- A relative importance parameter λ , rescaled by the number of repetitions of the cross-prediction task, is added.

3 Experiments and results

For comparability, we used the classification task of the whole time series as proposed by Yue et al. (2021) as a downstream task. Why not other downstream tasks?. This involved fitting an SVM classifier to the instance-level representation of the time series that had been derived by max-pooling over all individual timestamps. This allowed the classes of each instance to be predicted, and the performance was measured in terms of accuracy.

To perform experiments, we used the subset of 12 UEA datasets to benchmark against TS2Vec, on which Yue et al. (2021) reported SOTA performance in their paper. The convergence path of the adapted loss function for the StandWalkJump dataset as an example can be seen below.




Figure 2: Loss convergence of StandWalkJump dataset for $s=0.1$

The loss converges fairly quickly for the model without integrating the AR model. After 60 epochs, the loss jumps up since the additional objective of the AR model was added. The combined objective then converges to a similar low value as the initial model, although the convergence is not as smooth as without the combined objective.

Since our implementation was done in a slightly different Python environment, we compared TS2VecAR to the replicated results of TS2Vec with the default settings as specified in their repository. The results of the experiments¹ are shown below, with different specifications of the share of iterations which did not include the autoregressive temporal contrasting component.

Dataset	AR ($s=0$)	AR ($s=0.1$)	AR ($s=0.2$)	Replicated	Type
SelfRegulationSCP2	0.544	0.550	0.550	0.556	EEG
StandWalkJump	0.533	0.467	0.400	0.467	ECG
SpokenArabicDigits	0.967	0.959	0.977	0.989	Speech
DuckDuckGeese	0.460	0.460	0.540	0.520	Audio
ArticulatoryWordRecognition	0.987	0.970	0.977	0.977	Motion
CharacterTrajectories	0.991	0.993	0.994	0.992	Motion
EigenWorms	0.786	0.756	0.809	0.863	Motion
PenDigits	0.988	0.986	0.990	0.989	Motion
Handwriting	0.556	0.551	0.548	0.531	HAR
NATOPS	0.911	0.839	0.878	0.939	HAR
RacketSports	0.888	0.888	0.895	0.855	HAR
UWaveGestureLibrary	0.919	0.925	0.919	0.906	HAR
Mean (All datasets)	0.794	0.779	0.790	0.799	
Mean (HAR datasets)	0.819	0.801	0.810	0.808	

In 8 out of 12 datasets TS2Vec is outperformed by at least one specification of TS2VecAR. However, on average, TS2Vec still has the highest accuracy, partly due to the very low performance of the $s = 0$ and $s = 1$ specifications in the DuckDuckGeese and Eigenworms datasets, and the $s = 0.2$ specification in the StandWalkJump and EigenWorms and NATOPS datasets.

¹As hyperparameters we chose $\lambda = 5$ since the ATC loss is smaller than the HC loss and $m = 5$ to allow for sufficient cross prediction tasks in each iteration. All models were trained for the 600 iterations as specified default in TS2Vec.

Considering only the subsample of Human Activity Recognition (HAR) datasets, TS2VecAR outperforms TS2Vec on three out of four datasets in all specifications. On average, the improvement is 1.1% when comparing the $s = 0$ specification to TS2Vec. However, for some datasets better individual performances can be achieved by including the autoregressive temporal contrasting after some initial training using only hierarchical contrasting. This suggests that it may be beneficial to first learn a better representation before using it for the cross-prediction task. However, the average accuracy for these specifications is compromised by the low performance on the NATOPS dataset.

4 Conclusion

TS2Vec has shown very strong results in learning robust latent representations of a time series via contextual contrasting, which can be used for various downstream tasks such as classification. However, the latent representation does not capture information about future time stamps, which would allow to learn more about the structure of the time series. Thus, TS2Vec can be extended by adding an autoregressive component in the learning process, which performs a cross-prediction task. This has shown better results for several datasets, where this study has in particular shown improvements for human activity recognition tasks.

References

- Bachman, P., Hjelm, R. D., and Buchwalter, W. (2019). Learning representations by maximizing mutual information across views. *CoRR*, abs/1906.00910.
- Chen, T., Kornblith, S., Norouzi, M., and Hinton, G. E. (2020). A simple framework for contrastive learning of visual representations. *CoRR*, abs/2002.05709.
- Ching, T., Himmelstein, D. S., Beaulieu-Jones, B. K., Kalinin, A. A., Do, B. T., Way, G. P., Ferrero, E., Agapow, P.-M., Zietz, M., Hoffman, M. M., Xie, W., Rosen, G. L., Lengerich, B. J., Israeli, J., Lanchantin, J., Woloszynek, S., Carpenter, A. E., Shrikumar, A., Xu, J., Cofer, E. M., Lavender, C. A., Turaga, S. C., Alexandari, A. M., Lu, Z., Harris, D. J., DeCaprio, D., Qi, Y., Kundaje, A., Peng, Y., Wiley, L. K., Segler, M. H., Boca, S. M., Swamidass, S. J., Huang, A., Gitter, A., and Greene, C. S. (2018). Opportunities and obstacles for deep learning in biology and medicine. *bioRxiv*.
- Eldele, E., Ragab, M., Chen, Z., Wu, M., Kwoh, C.-K., Li, X., and Guan, C. (2022). Self-supervised contrastive representation learning for semi-supervised time-series classification.
- Franceschi, J.-Y., Dieuleveut, A., and Jaggi, M. (2019). Unsupervised scalable representation learning for multivariate time series.
- Mohsenvand, M. N., Izadi, M. R., and Maes, P. (2020). Contrastive representation learning for electroencephalogram classification. In Alsentzer, E., McDermott, M. B. A., Falck, F., Sarkar, S. K., Roy, S., and Hyland, S. L., editors, *Proceedings of the Machine Learning for Health NeurIPS Workshop*, volume 136 of *Proceedings of Machine Learning Research*, pages 238–253. PMLR.
- Oord, A. v. d., Li, Y., and Vinyals, O. (2018). Representation learning with contrastive predictive coding.
- Tonekaboni, S., Eytan, D., and Goldenberg, A. (2021). Unsupervised representation learning for time series with temporal neighborhood coding.
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, L., and Polosukhin, I. (2017). Attention is all you need.
- Yue, Z., Wang, Y., Duan, J., Yang, T., Huang, C., Tong, Y., and Xu, B. (2021). TS2Vec: Towards universal representation of time series.